

# Transformers

Dr Mehreen Alam


<https://jalammar.github.io/illustrated-transformer/>

<https://machinelearningmastery.com/a-gentle-introduction-to-positional-encoding-in-transformer-models-part-1/>

# Challenges of RNNs


- Long range dependencies
- Parallelization

## English German Translation quality



BLEU scores (higher is better) of single models on the standard WMT newstest2014 English to German translation benchmark.



## English French Translation Quality

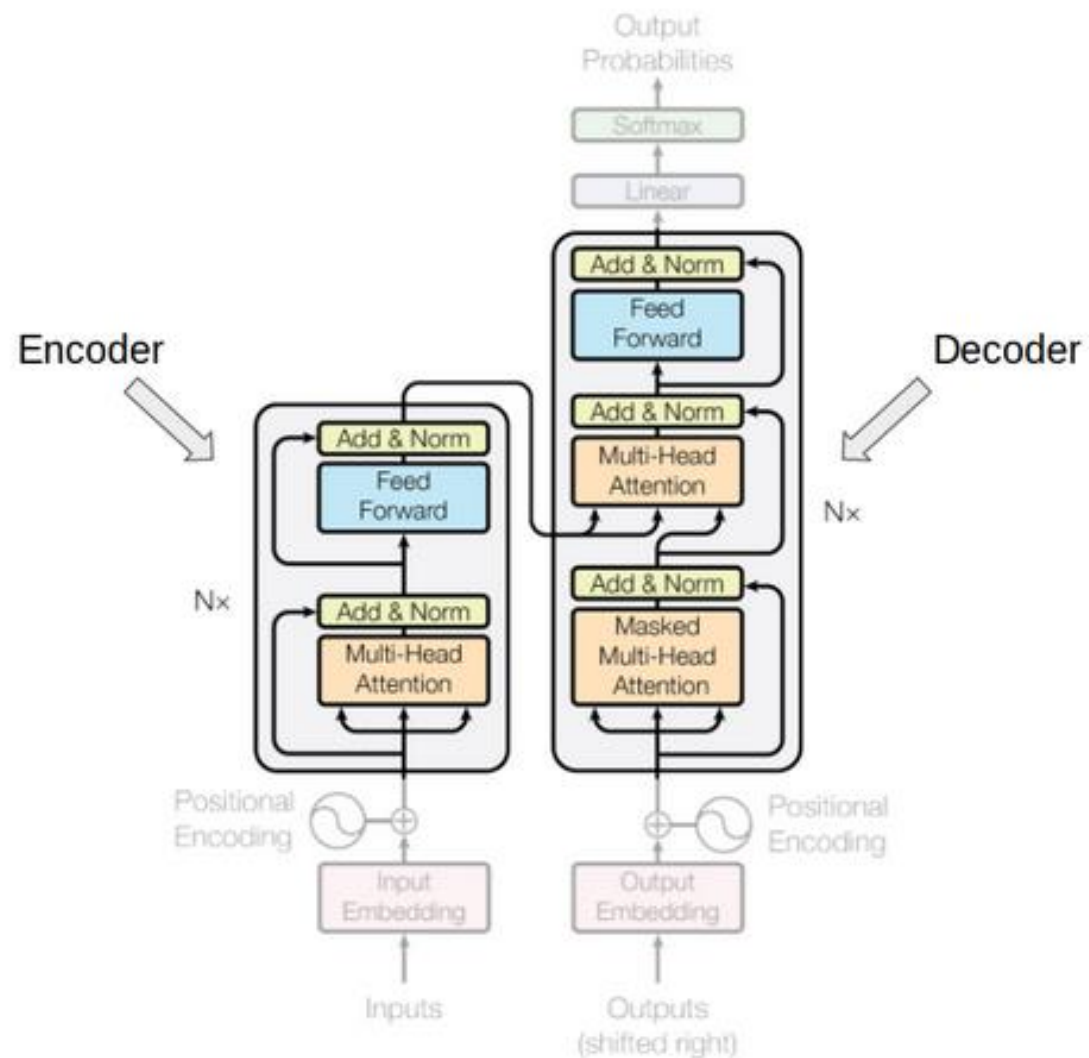
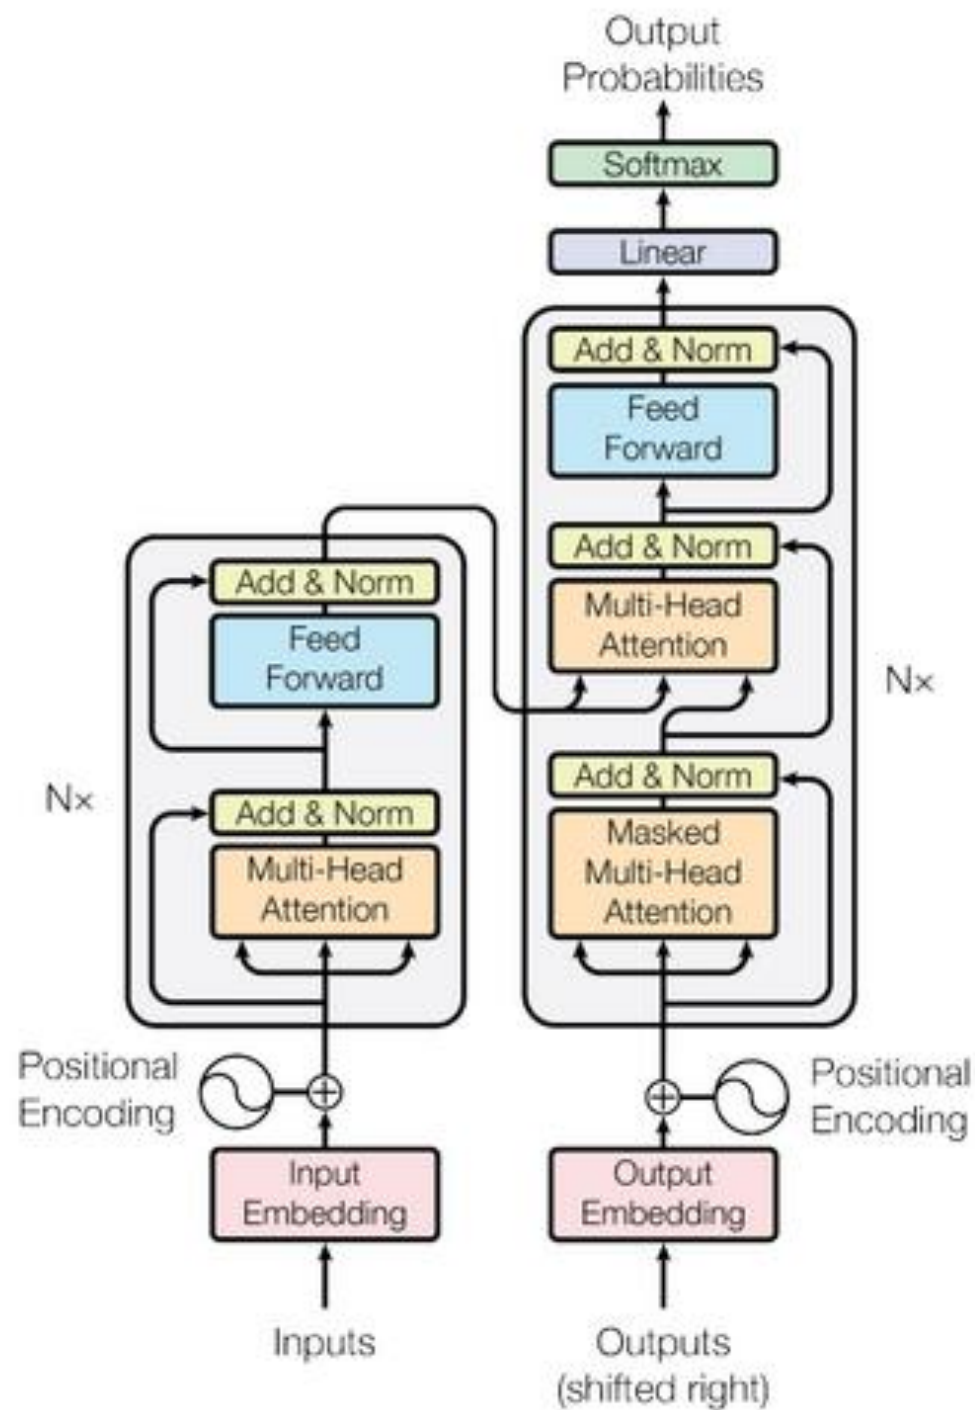


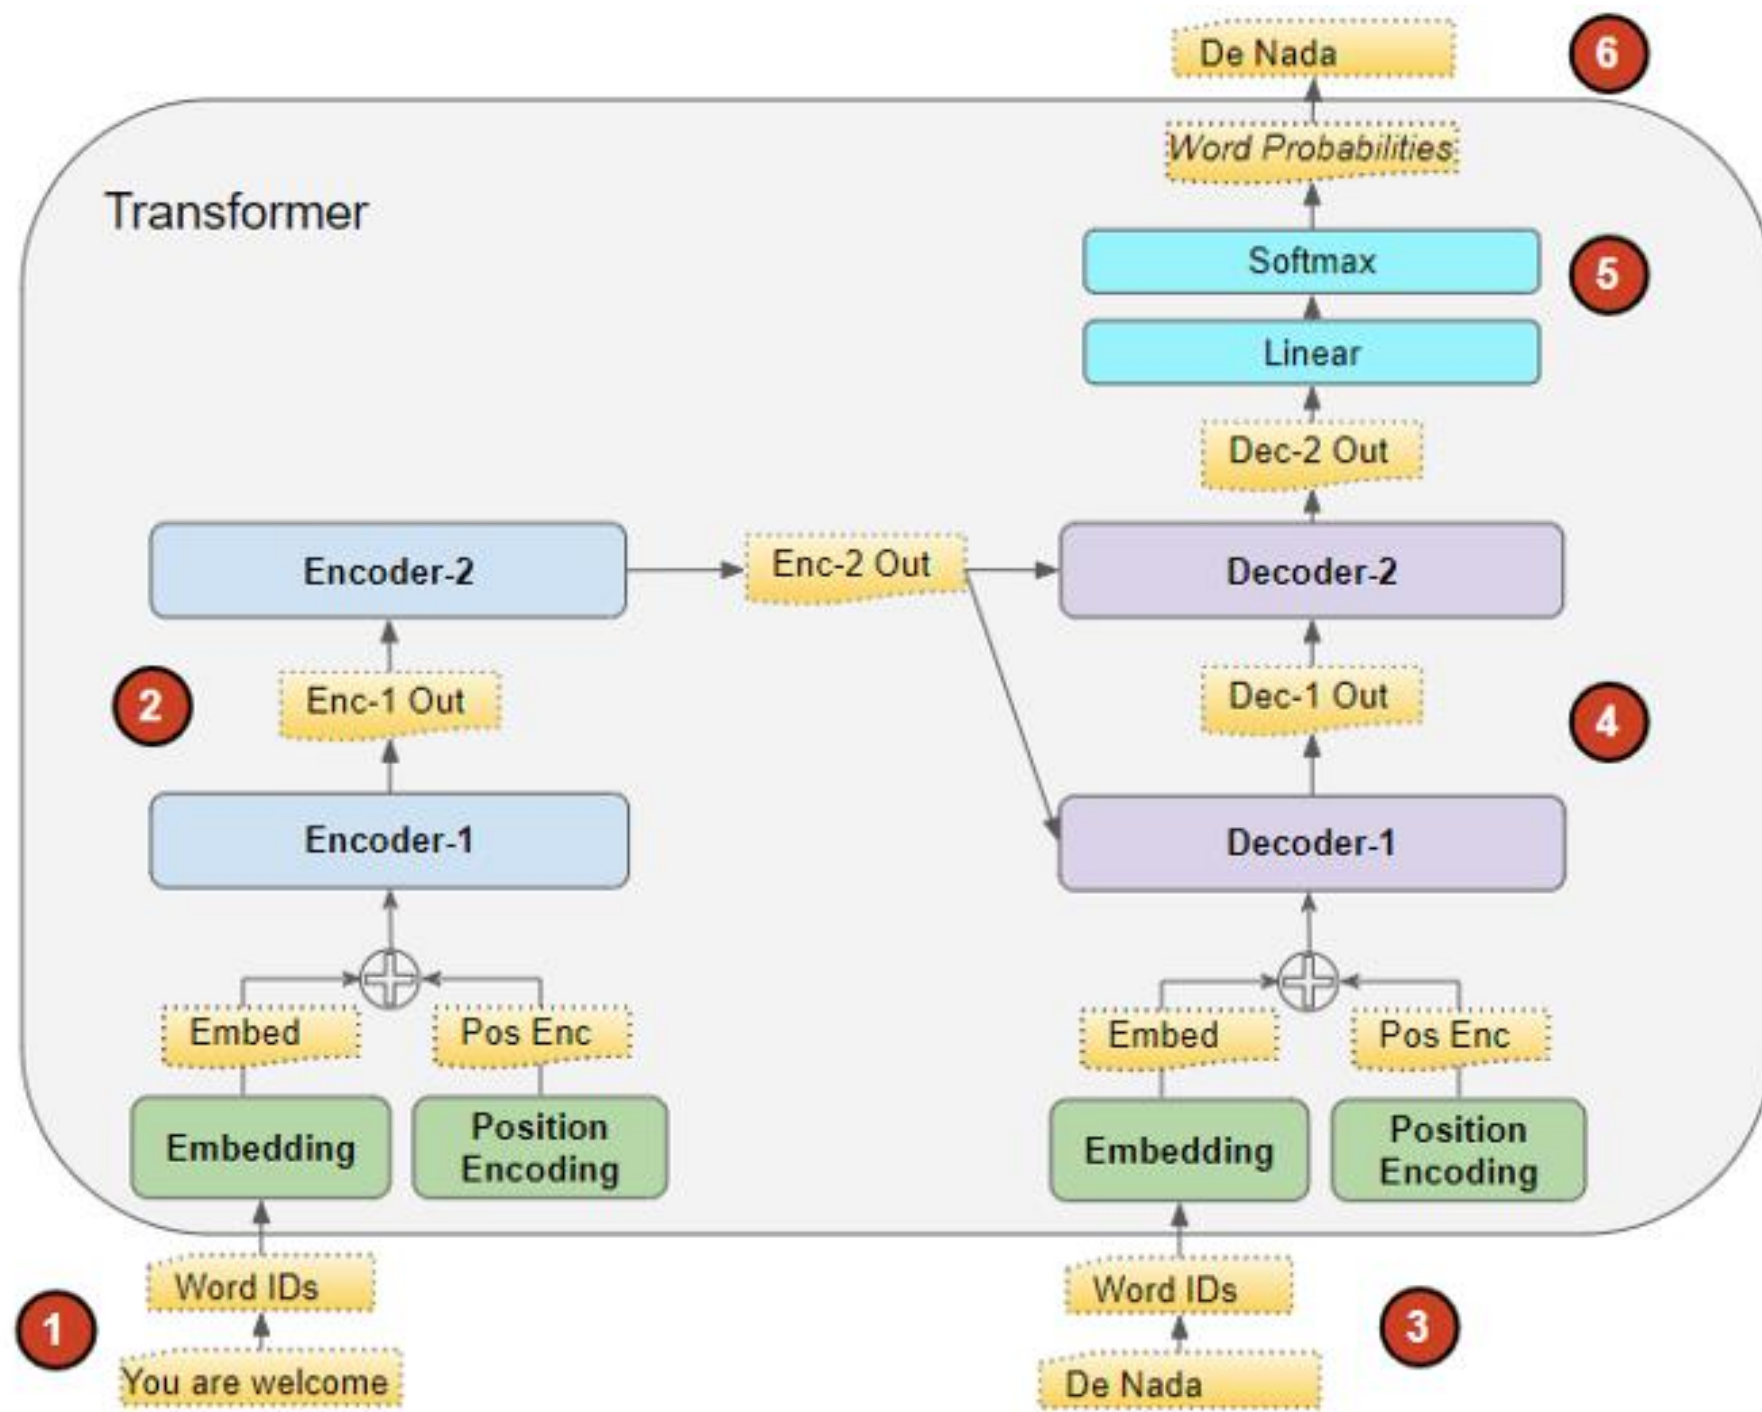
BLEU scores (higher is better) of single models on the standard WMT newstest2014 English to French translation benchmark.

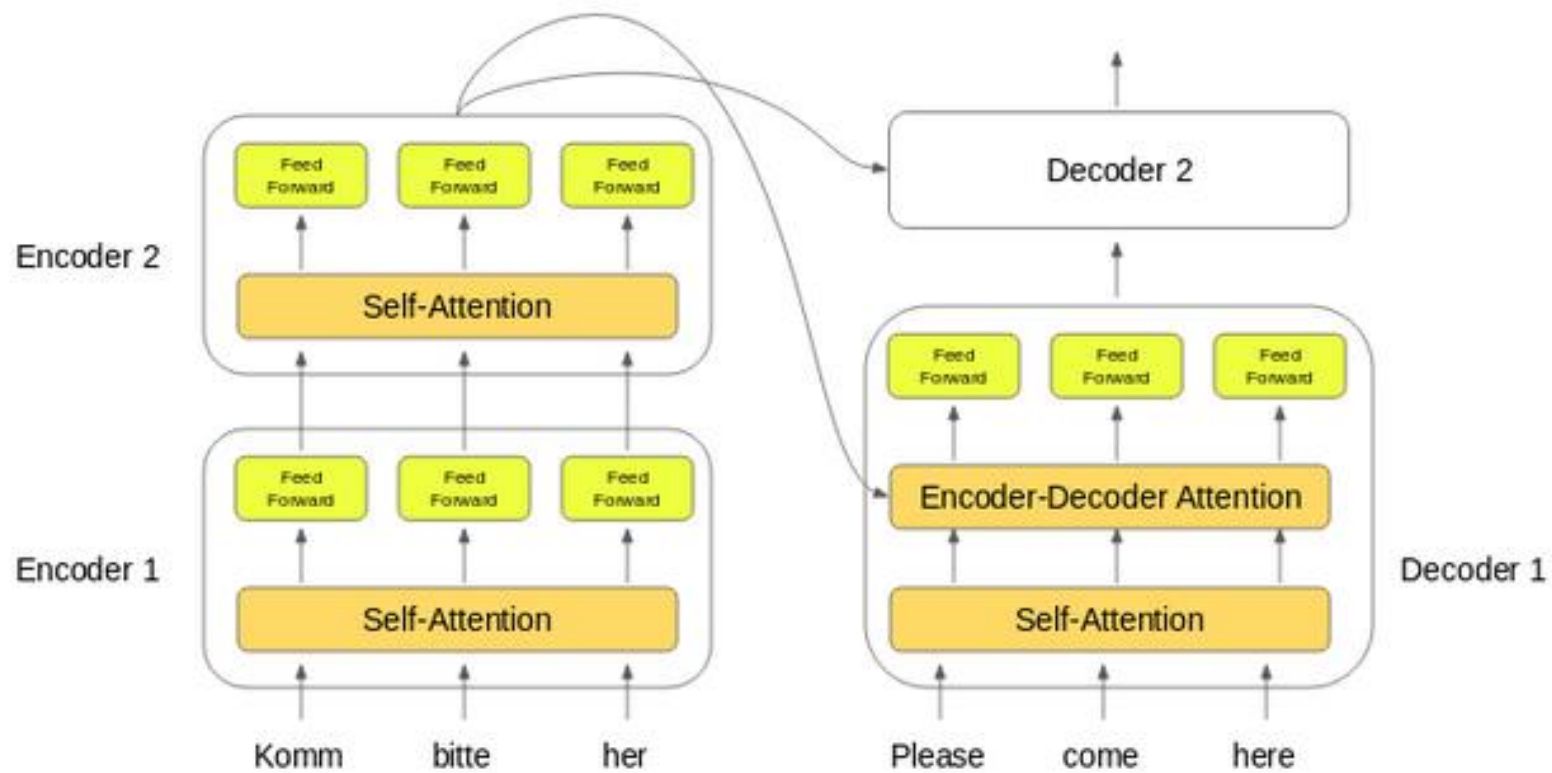
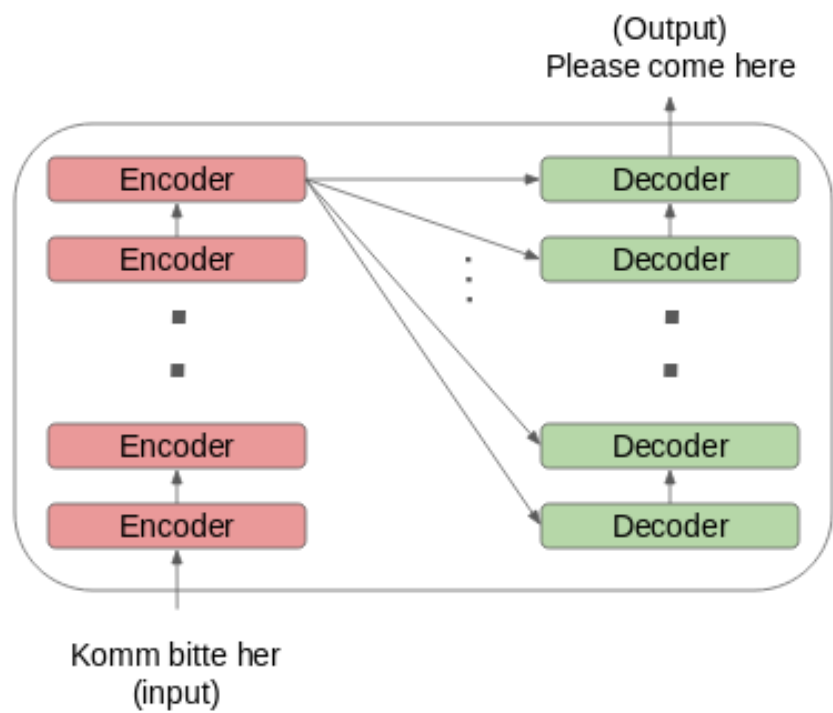
*The animal didn't cross the street because it was too tired.*  
*L'animal n'a pas traversé la rue parce qu'il était trop fatigué.*

*The animal didn't cross the street because it was too wide.*  
*L'animal n'a pas traversé la rue parce qu'elle était trop large.*









# Self-Attention

- attention mechanism relating different positions of a single sequence in order to compute a representation of the sequence.
- intra-attention





# Calculating self-attention

1. create three vectors from each of the encoder's input vectors:
  1. Query Vector
  2. Key Vector
  3. Value Vector
2. calculate self-attention for every word in the input sequence
  1. Example for "Action gets results."

Word	q vector	k vector	v vector
Action	$q_1$	$k_1$	$v_1$
gets		$k_2$	$v_2$
results		$k_3$	$v_3$

Word	q vector	k vector	v vector	score
Action	$q_1$	$k_1$	$v_1$	$q_1 \cdot k_1$
gets		$k_2$	$v_2$	$q_1 \cdot k_2$
results		$k_3$	$v_3$	$q_1 \cdot k_3$

Word	q vector	k vector	v vector	score	score / 8
Action	$q_1$	$k_1$	$v_1$	$q_1 \cdot k_1$	$q_1 \cdot k_1 / 8$
gets		$k_2$	$v_2$	$q_1 \cdot k_2$	$q_1 \cdot k_2 / 8$
results		$k_3$	$v_3$	$q_1 \cdot k_3$	$q_1 \cdot k_3 / 8$

Word	q vector	k vector	v vector	score	score / 8	Softmax
Action	$q_1$	$k_1$	$v_1$	$q_1 \cdot k_1$	$q_1 \cdot k_1 / 8$	$x_{11}$
gets		$k_2$	$v_2$	$q_1 \cdot k_2$	$q_1 \cdot k_2 / 8$	$x_{12}$
results		$k_3$	$v_3$	$q_1 \cdot k_3$	$q_1 \cdot k_3 / 8$	$x_{13}$

Word	q vector	k vector	v vector	score	score / 8	Softmax	Softmax * v	Sum
Action	$q_1$	$k_1$	$v_1$	$q_1 \cdot k_1$	$q_1 \cdot k_1 / 8$	$x_{11}$	$x_{11} * v_1$	$z_1$
gets		$k_2$	$v_2$	$q_1 \cdot k_2$	$q_1 \cdot k_2 / 8$	$x_{12}$	$x_{12} * v_2$	
results		$k_3$	$v_3$	$q_1 \cdot k_3$	$q_1 \cdot k_3 / 8$	$x_{13}$	$x_{13} * v_3$	

Word	q vector	k vector	v vector	score	score / 8	Softmax	Softmax * v	Sum <sup>#</sup>
Action		$k_1$	$v_1$	$q_2 \cdot k_1$	$q_2 \cdot k_1 / 8$	$x_{21}$	$x_{21} * v_1$	
gets	$q_2$	$k_2$	$v_2$	$q_2 \cdot k_2$	$q_2 \cdot k_2 / 8$	$x_{22}$	$x_{22} * v_2$	$z_2$
results		$k_3$	$v_3$	$q_2 \cdot k_3$	$q_2 \cdot k_3 / 8$	$x_{23}$	$x_{23} * v_3$	

Word	q vector	k vector	v vector	score	score / 8	Softmax	Softmax * v	Sum <sup>#</sup>
Action		$k_1$	$v_1$	$q_3 \cdot k_1$	$q_3 \cdot k_1 / 8$	$x_{31}$	$x_{31} * v_1$	
gets		$k_2$	$v_2$	$q_3 \cdot k_2$	$q_3 \cdot k_2 / 8$	$x_{32}$	$x_{32} * v_2$	
results	$q_3$	$k_3$	$v_3$	$q_3 \cdot k_3$	$q_3 \cdot k_3 / 8$	$x_{33}$	$x_{33} * v_3$	$z_3$

# Recap of Attention from Seq2Seq

## Attention: in equations

- We have encoder hidden states  $h_1, \dots, h_N \in \mathbb{R}^h$
- On timestep  $t$ , we have decoder hidden state  $s_t \in \mathbb{R}^h$
- We get the attention scores  $e^t$  for this step:

$$e^t = [s_t^T h_1, \dots, s_t^T h_N] \in \mathbb{R}^N$$

- We take softmax to get the attention distribution  $\alpha^t$  for this step (this is a probability distribution and sums to 1)

$$\alpha^t = \text{softmax}(e^t) \in \mathbb{R}^N$$


- We use  $\alpha^t$  to take a weighted sum of the encoder hidden states to get the attention output  $a_t$

$$a_t = \sum_{i=1}^N \alpha_i^t h_i \in \mathbb{R}^h$$

- Finally we concatenate the attention output  $a_t$  with the decoder hidden state  $s_t$  and proceed as in the non-attention seq2seq model


$$[a_t; s_t] \in \mathbb{R}^{2h}$$

## encoder-to-sequence with attention




Decoder RNN

# Multi-headed Attention



Multi-Head Attention


1) Concatenate all the attention heads



2) Multiply with a weight matrix  $W^O$  that was trained jointly with the model

$\times$

3) The result would be the  $Z$  matrix that captures information from all the attention heads. We can send this forward to the FFNN



1) This is our input sentence\*





2) We embed each word\*

3) Split into 8 heads.  
We multiply  $X$  or  $R$  with weight matrices

4) Calculate attention using the resulting  $Q/K/V$  matrices

5) Concatenate the resulting  $Z$  matrices, then multiply with weight matrix  $W^O$  to produce the output of the layer



Thinking  
Machines



$W^O$






\* In all encoders other than #0,  
we don't need embedding.  
We start directly with the output  
of the encoder right below this one



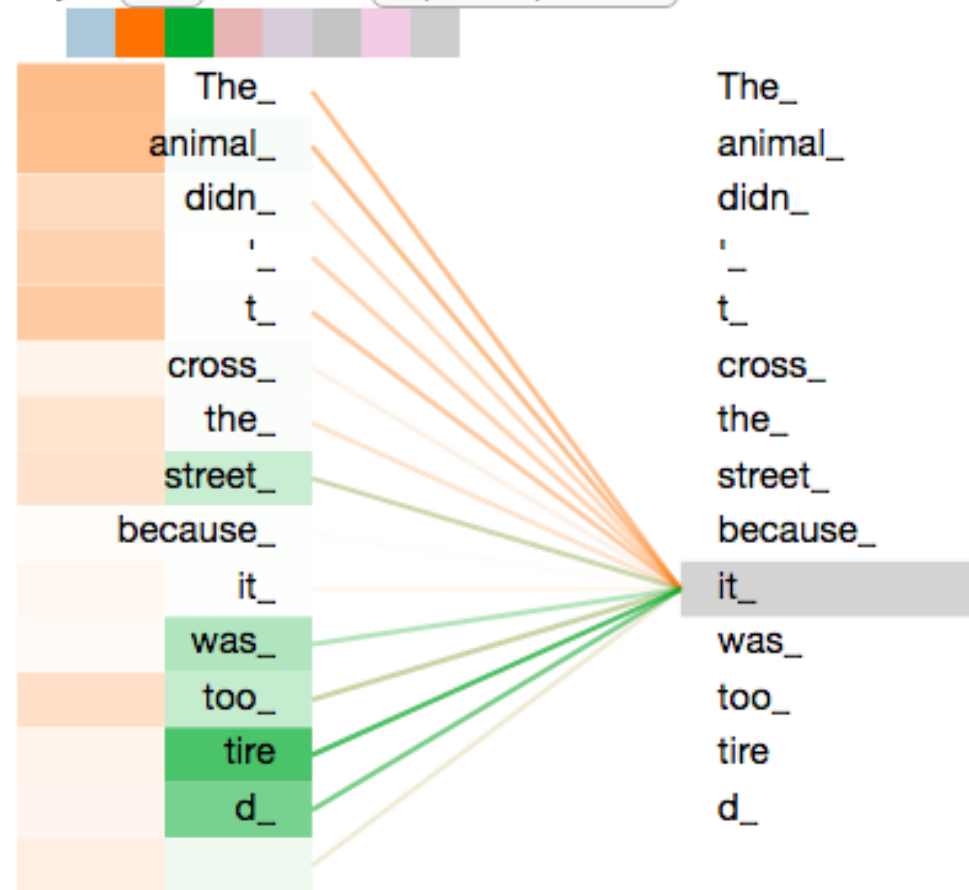
...

...

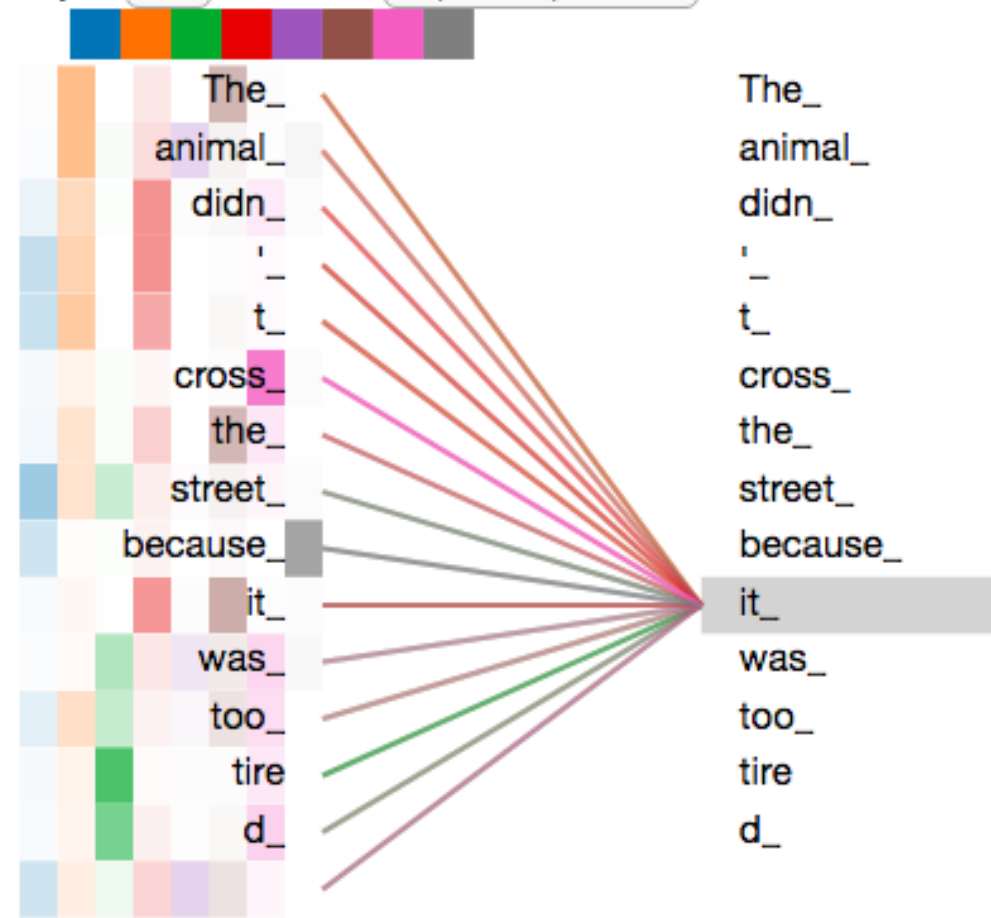
...



Layer: 5 Attention: Input - Input





Layer: 5 Attention: Input - Input



# Positional Encoding and Embedding

POSITIONAL  
ENCODING

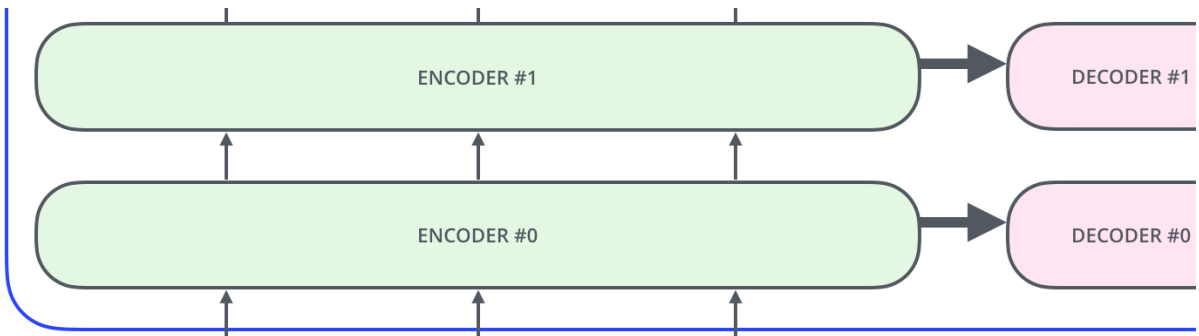





+




+

+

EMBEDDINGS



EMBEDDING  
WITH TIME  
SIGNAL




=

=

=

POSITIONAL  
ENCODING






+

+

+

EMBEDDINGS



INPUT

Je

suis


étudiant

INPUT

Je

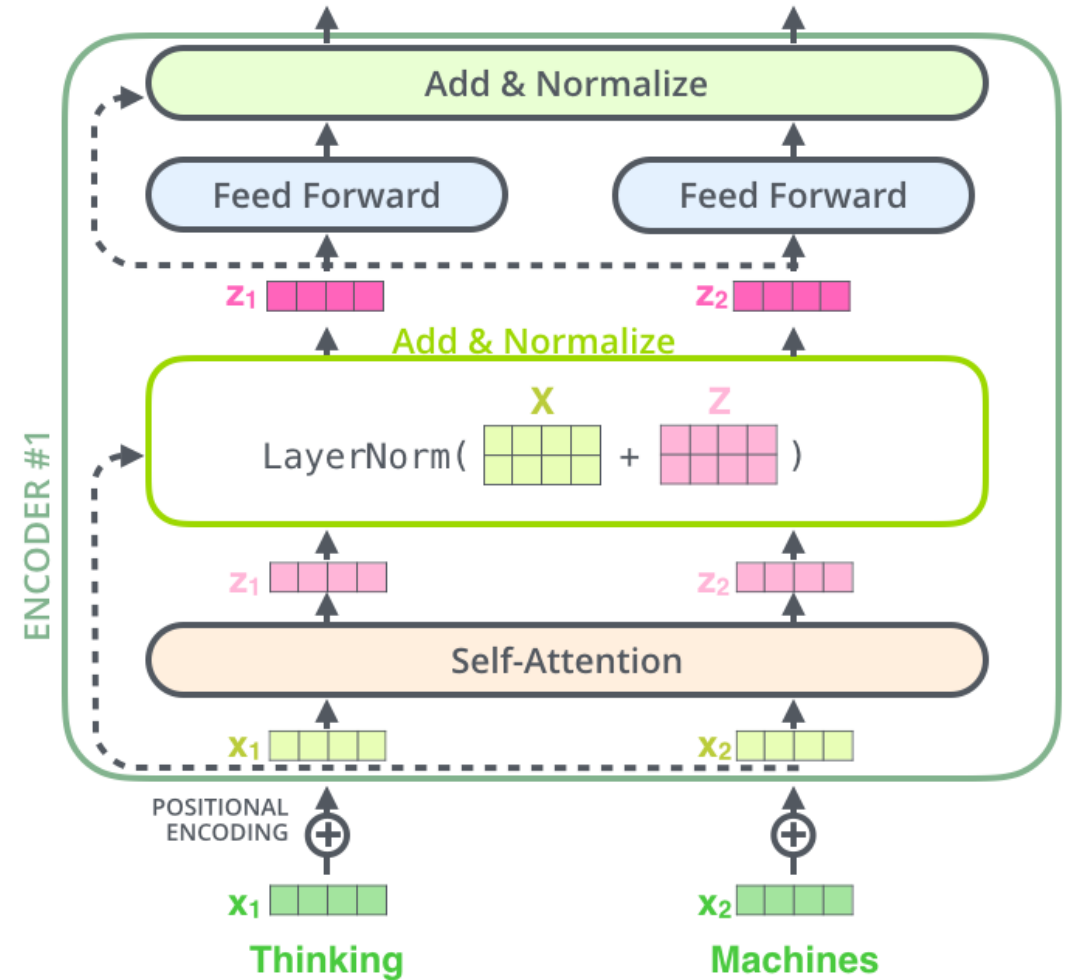
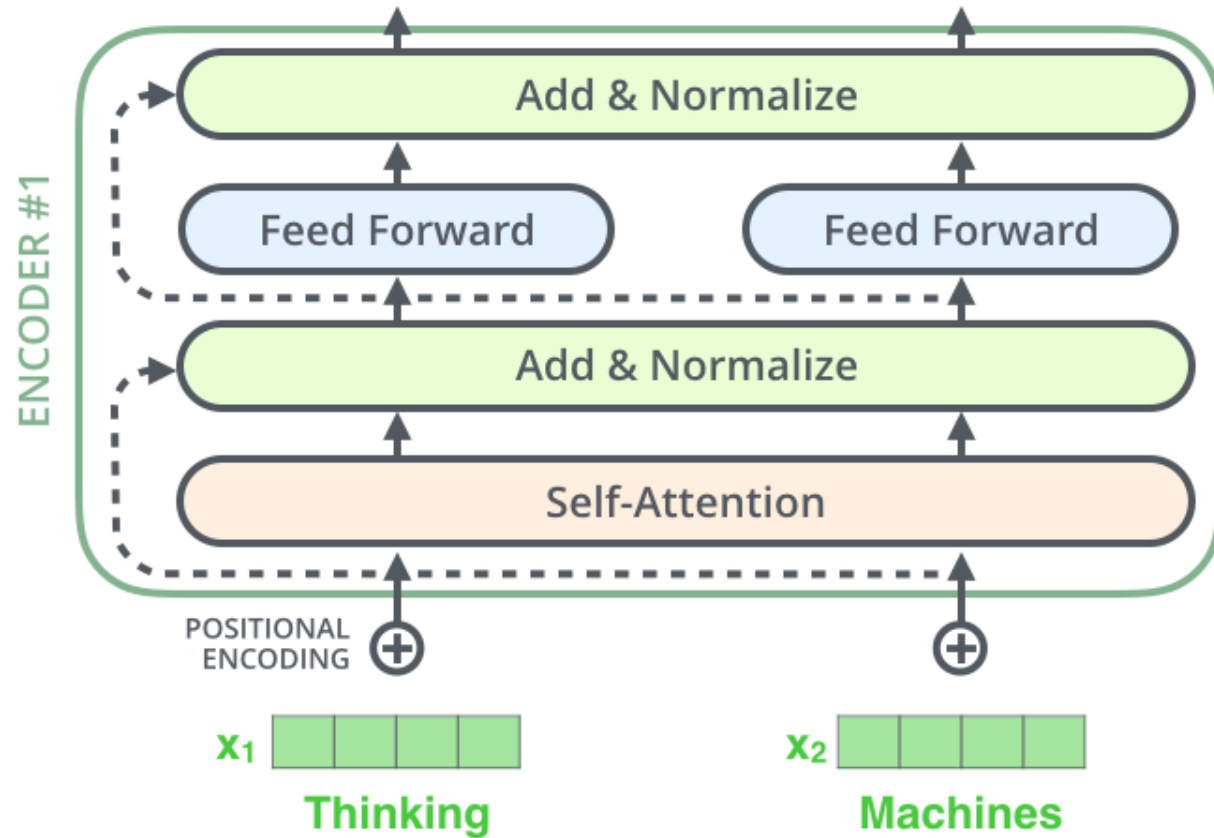
suis

étudiant





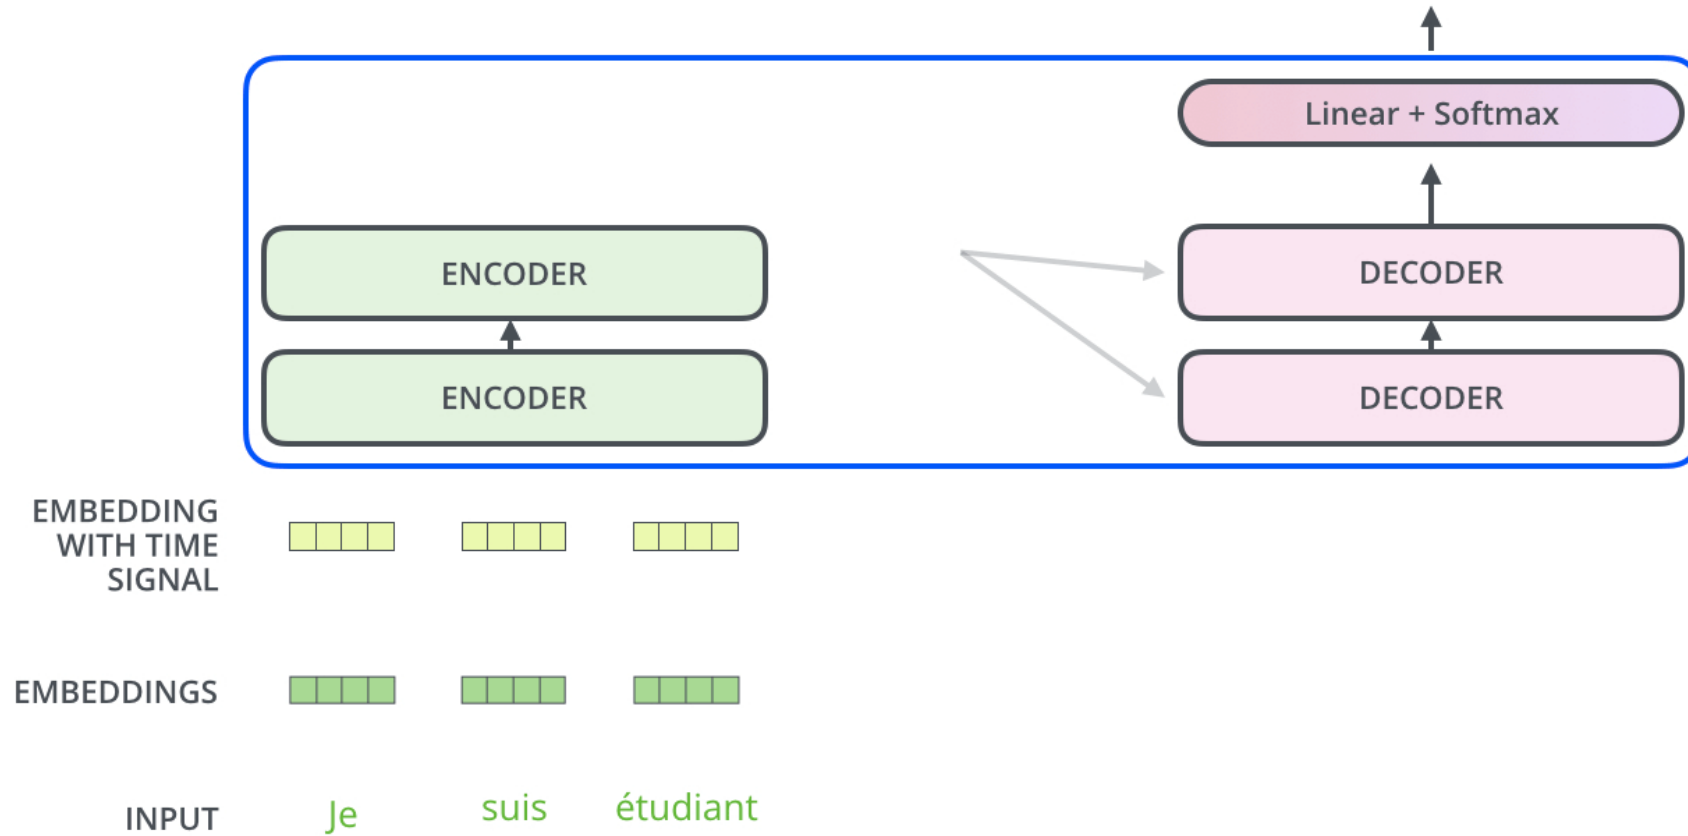
# Residuals



# Decoder

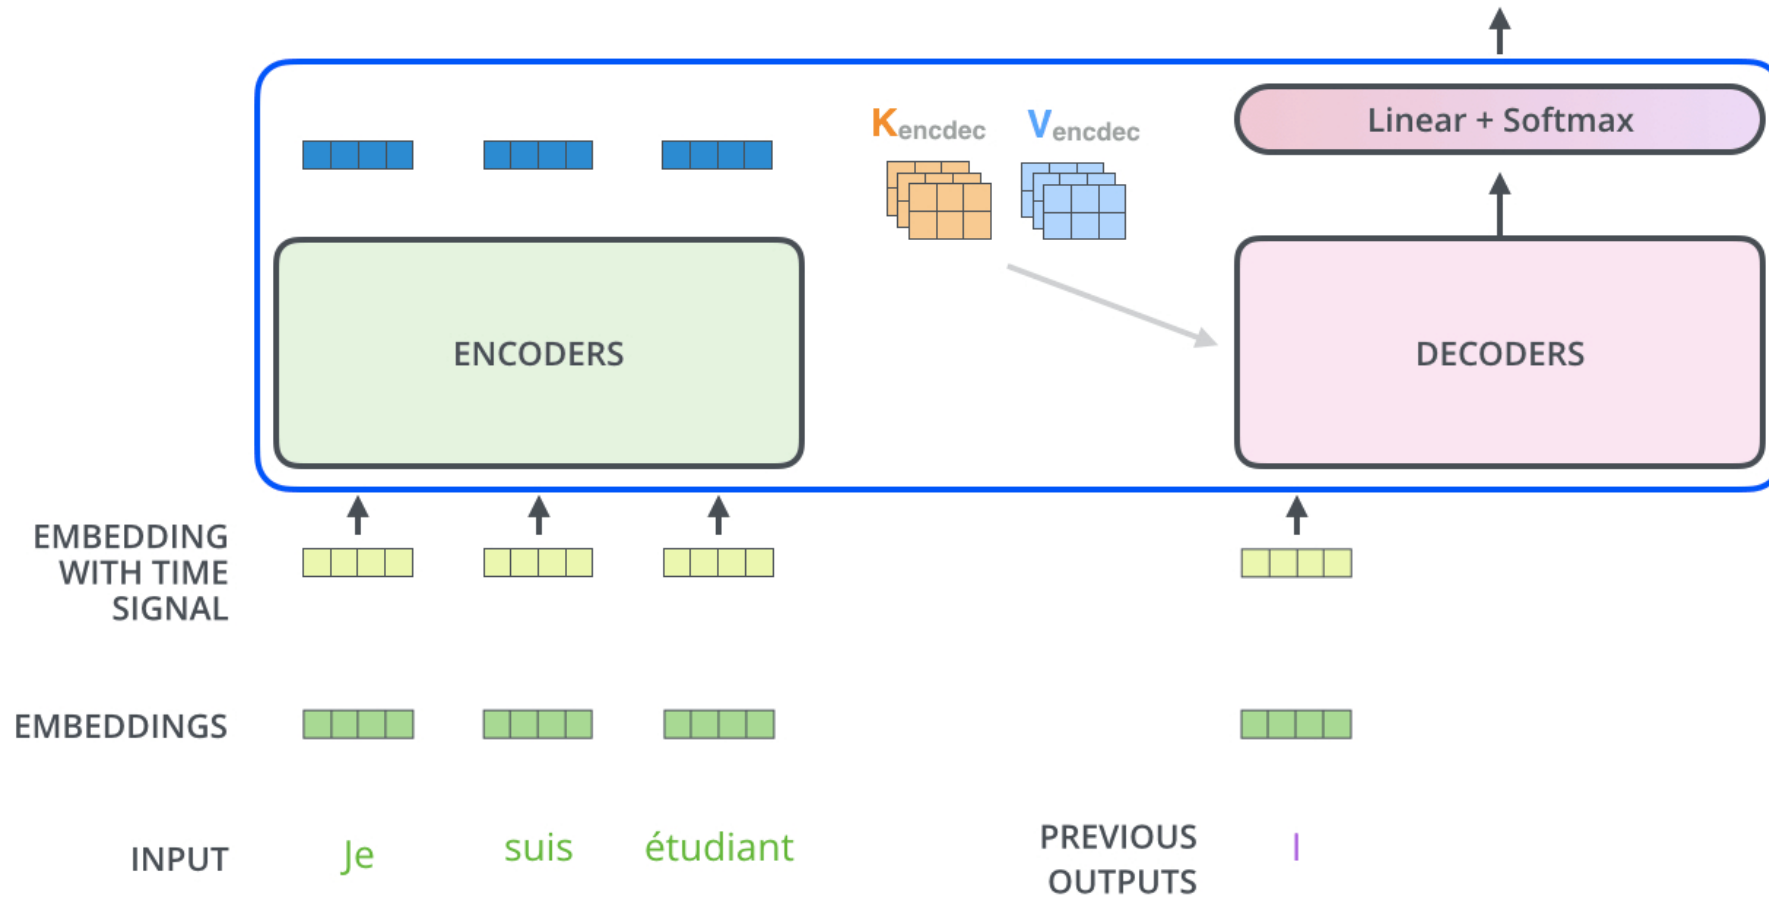
Decoding time step: ① 2 3 4 5 6

OUTPUT



Decoding time step: 1 2 3 4 5 6

OUTPUT |



Which word in our vocabulary  
is associated with this index?

Get the index of the cell  
with the highest value  
(**argmax**)

am

5

log\_probs



Softmax

logits



Linear

Decoder stack output



# Training Details

Output Vocabulary

WORD	a	am	I	thanks	student	<eos>
INDEX	0	1	2	3	4	5

One-hot encoding of the word "am"



Untrained Model Output



Correct and desired output



a

am

I

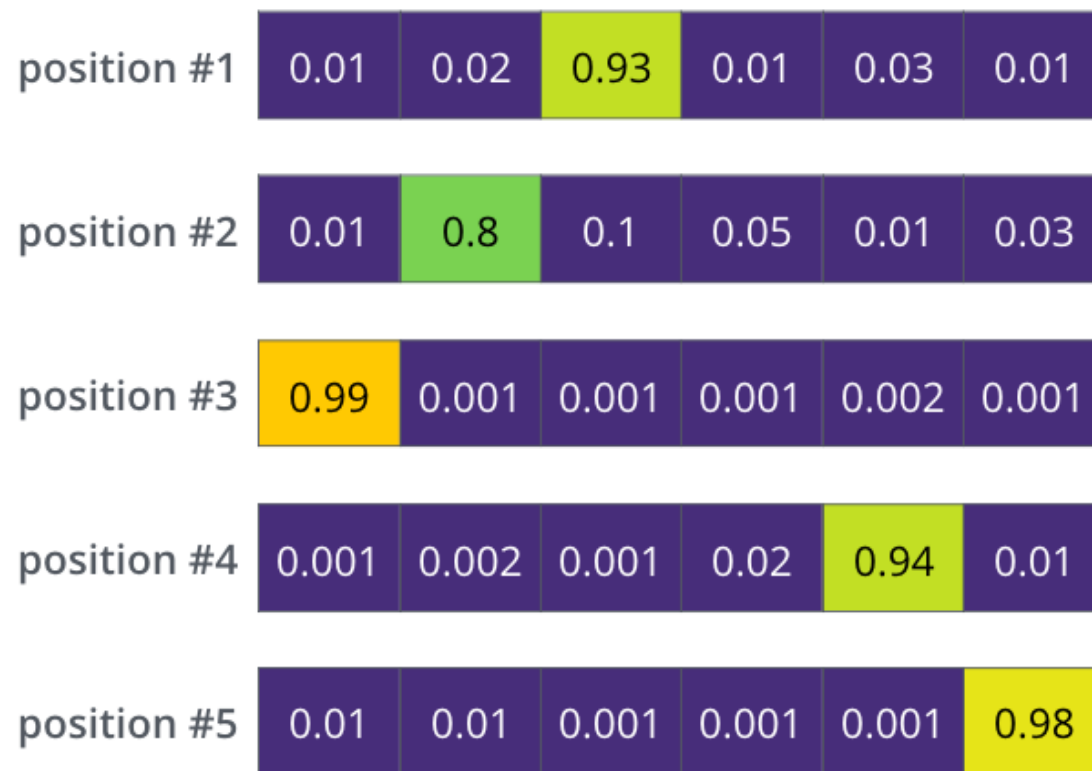
thanks

student

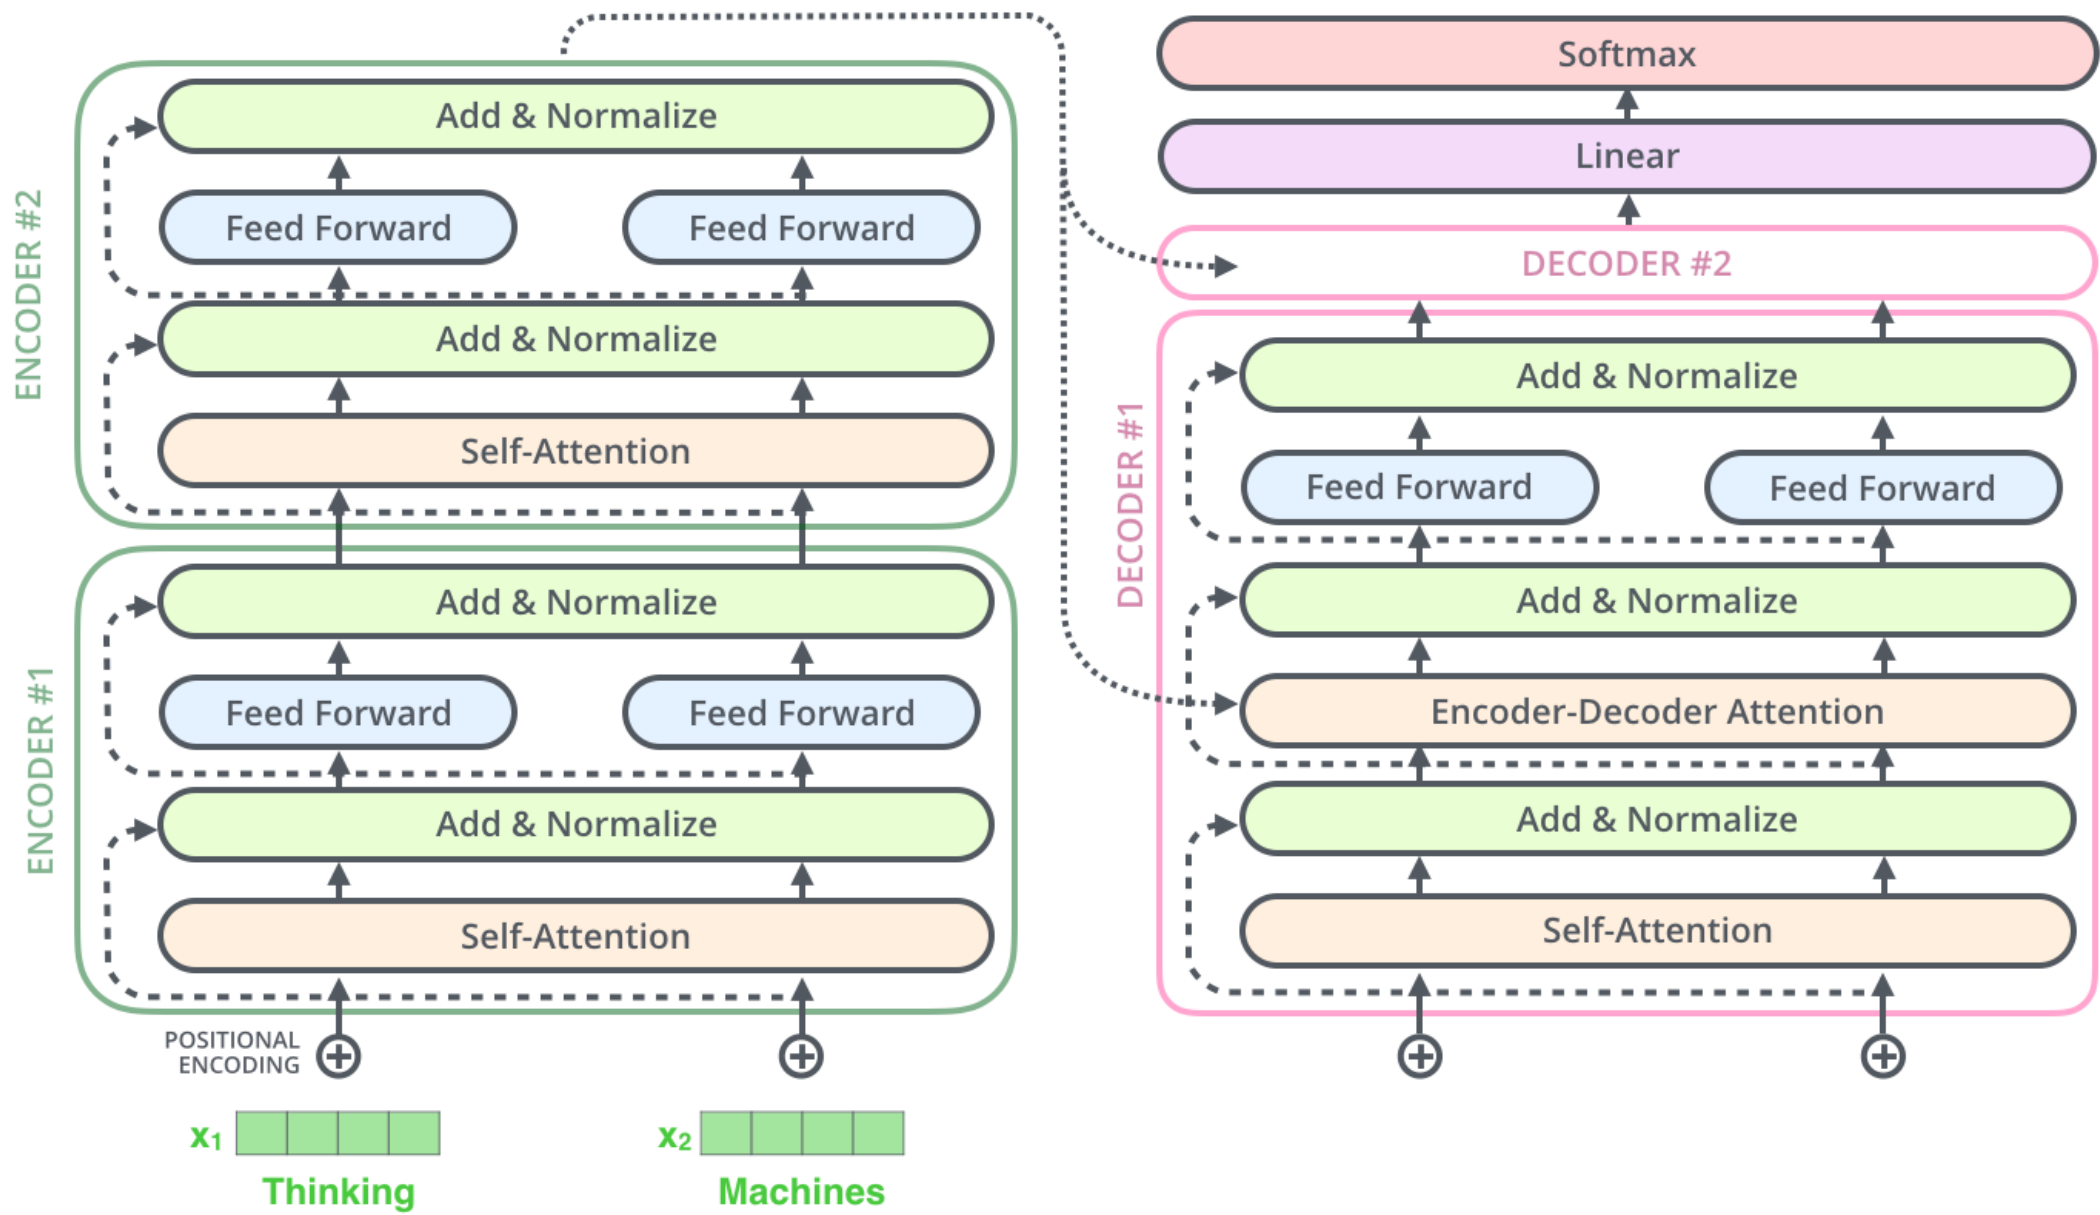
<eos>

## Trained Model Outputs

Output Vocabulary:    a        am        I        thanks    student    <eos>



a        am        I        thanks    student    <eos>





# Limitations of Transformers

- Attention can only deal with fixed-length text strings:
  - The text has to be split into a certain number of segments before being fed into the system as input.
- This chunking of text causes context fragmentation:
  - E.g., if a sentence is split from the middle, then a significant amount of context is lost.
  - In other words, the text is split without respecting the sentence or any other semantic boundary

# Transformers and Self-Attention: Mathematical Example

- Input Embedding and Positional Encoding
- Multi-Head Self-Attention (Encoder)
- Feed Forward Network (Encoder)
- Masked Multi-Head Attention (Decoder)Encoder-Decoder Attention
- Final Linear Layer and Softmax

# Transformers and Self-Attention: Mathematical Example

- For the given sentence
  - Input: "The animal didn't cross the street because it was too tired."
  - Output: "L'animal n'a pas traversé la rue parce qu'il était trop fatigué."

# Step 1: Input Embedding and Positional Encoding

## a) Input Token Embedding

The input sentence, "The animal didn't cross the street because it was too tired.", is tokenized, and each token is converted into a vector of size  $d_{\text{model}} = 512$ .

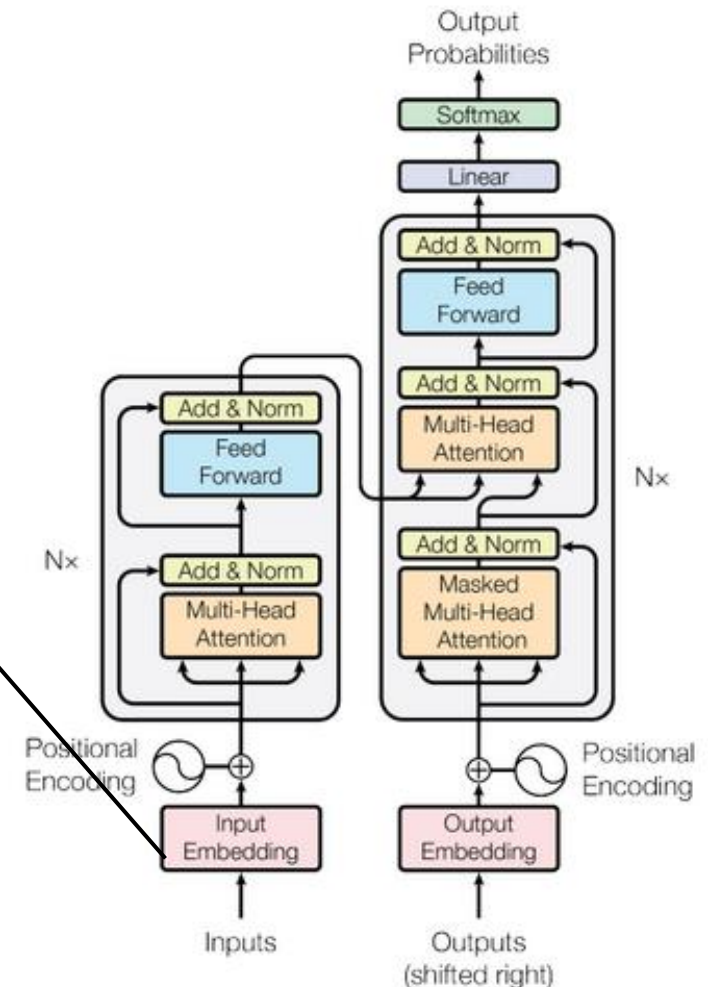
- **Tokens:** ["The", "animal", "didn't", "cross", "the", "street", "because", "it", "was", "too", "tired"]
- **Embedding Matrix Representation:** Each token is embedded into a 512-dimensional vector.

$$X = [x_1 \ x_2 \ \cdots \ x_n]^T, \quad X \in \mathbb{R}^{n \times d_{\text{model}}}$$

Where  $x_i$  is the embedding vector for the  $i$ -th token, and  $n$  is the number of tokens.

So, if the sentence has 10 tokens, you will have 10 vectors, each of size 512.

Each word or token in the sentence gets **its own 512-dimensional vector**.



# Step 1: Input Embedding and Positional Encoding

## b) Positional Encoding

Since Transformers don't have built-in information about the positions of words, we add **positional encodings** to the embeddings. These encodings are based on sine and cosine functions of different frequencies.

The positional encoding is defined as:

$$PE_{(pos, 2i)} = \sin\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right)$$
$$PE_{(pos, 2i+1)} = \cos\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right)$$

Where:

- $pos$  is the position of the word in the sentence.
- $2i$  and  $2i + 1$  refer to even and odd dimensions in the positional encoding vector.

For example, if  $d_{model} = 512$ , then:

- $PE_{(1,0)} = \sin\left(\frac{1}{10000^{0/512}}\right) = \sin(1)$
- $PE_{(1,1)} = \cos\left(\frac{1}{10000^{0/512}}\right) = \cos(1)$
- And so on for all positions and dimensions.

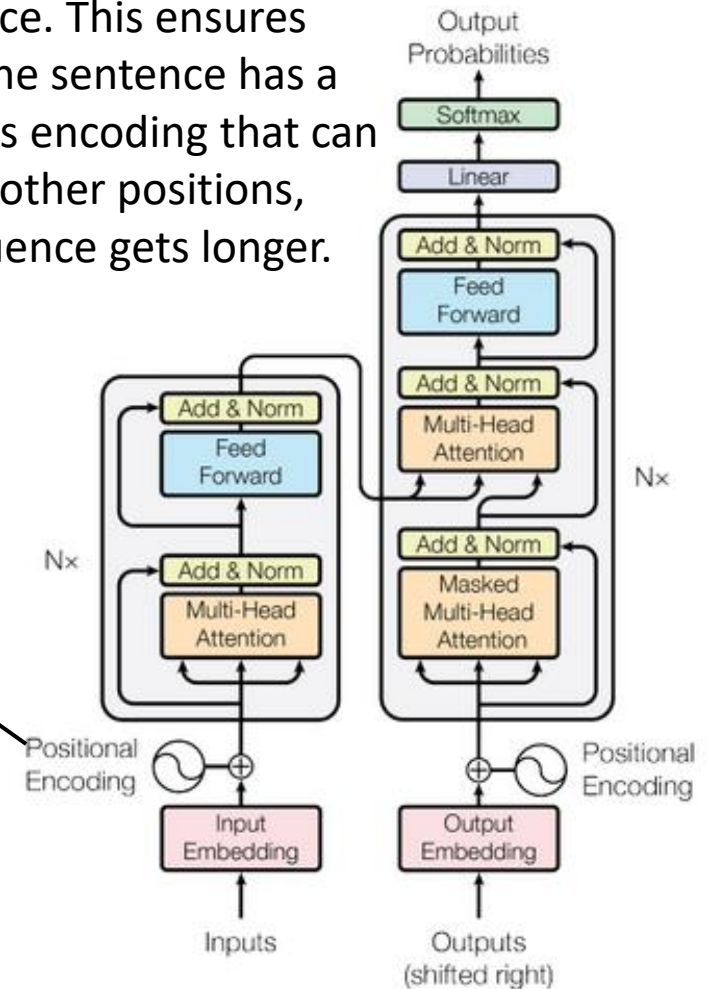
The final positional encoding matrix  $PE \in \mathbb{R}^{n \times d_{model}}$  is added to the input embedding matrix  $X$ :

$$X' = X + PE$$

This creates the final matrix  $X'$ , which is passed to the encoder.

By alternating between sine and cosine functions, the model can create a **rich and unique representation** for each position in the sequence. This ensures that each position in the sentence has a distinct and continuous encoding that can be distinguished from other positions, even as the input sequence gets longer.

The even-odd split allows the model to handle a wider range of positions effectively and gives it the ability to **distinguish between positions that are close together** (like adjacent words) and **positions that are far apart** (words at the beginning and end of a sentence).



## 2. Multi-Head Self-Attention (Encoder)

<https://aibyhand.substack.com/p/11-can-you-calculate-self-attention>

The self-attention mechanism enables each word in the sentence to attend to every other word, helping the model understand relationships between words.

### a) Queries, Keys, and Values

For each input embedding  $x_i$ , we project it into three different vectors:

- Query (Q):  $Q = X'W_Q$
- Key (K):  $K = X'W_K$
- Value (V):  $V = X'W_V$

Where  $W_Q, W_K, W_V \in \mathbb{R}^{d_{\text{model}} \times d_k}$  are the learned projection matrices, and  $d_k$  is typically  $d_{\text{model}}/h$  (e.g., for 8 heads,  $d_k = 512/8 = 64$ ).

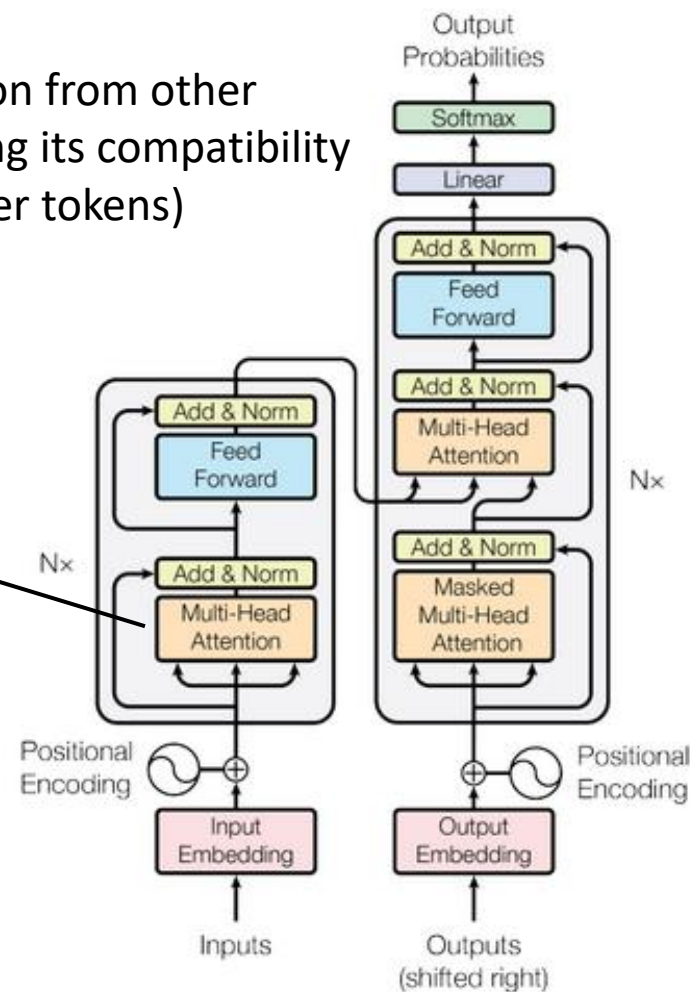
For  $X' \in \mathbb{R}^{n \times d_{\text{model}}}$ , the resulting matrices are:

$$Q, K, V \in \mathbb{R}^{n \times d_k}$$

The Value vector contains the word's representation and is used to compute the final weighted output after attention scores are applied. The Values are what get "returned" after the attention process decides which tokens to focus on.

The query is used to find relevant information from other tokens (words) in the sequence by computing its compatibility with Keys (which are representations of other tokens)

Each word in the sequence has a Key, which is compared to Queries from other tokens to compute attention scores. The Key holds information about a word and helps decide how important this word is when processing the current word (Query).



## 2. Multi-Head Self-Attention (Encoder)

<https://aibychand.substack.com/p/11-can-you-calculate-self-attention>

### b) Scaled Dot-Product Attention

For each word, we calculate the attention score between its **query** and every other word's **key** using a dot product:

$$\text{Attention Score} = QK^T$$

The attention score is then scaled by  $\frac{1}{\sqrt{d_k}}$  to prevent large values when the dot products are large:

$$\text{Scaled Score} = \frac{QK^T}{\sqrt{d_k}}$$

Next, we apply a **softmax** function to obtain the attention weights:

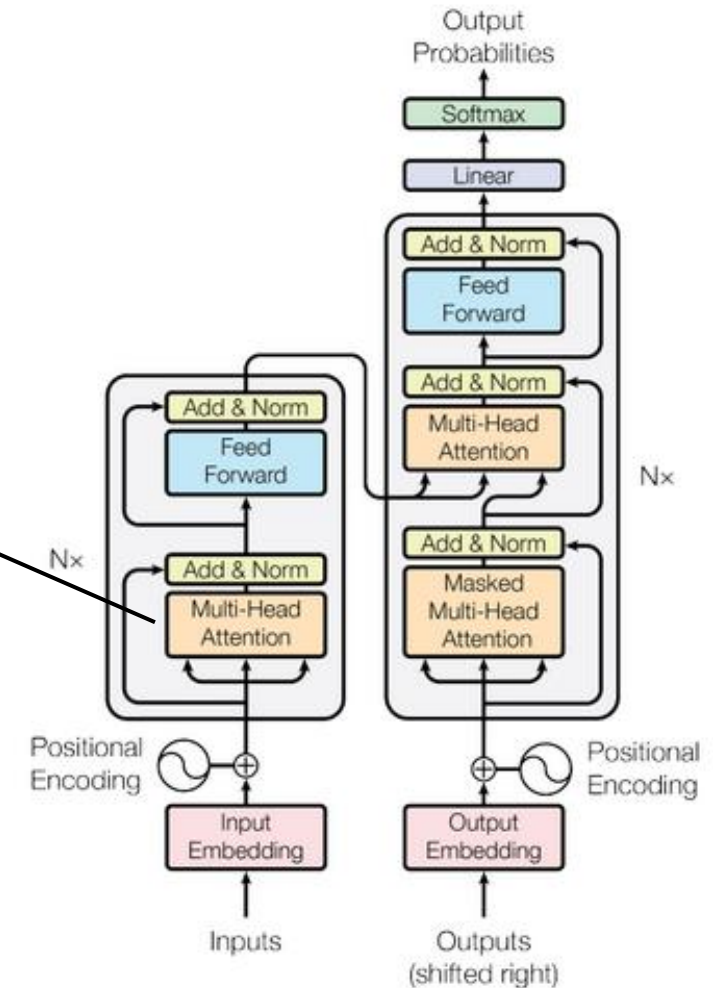
$$\text{Attention Weights} = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)$$

The attention weights are used to create a weighted sum of the **value** vectors:

$$\text{Attention Output} = \text{Attention Weights} \times V$$

This results in an output matrix:

$$\text{Attention Output} \in \mathbb{R}^{n \times d_k}$$





## 2. Multi-Head Self-Attention (Encoder)

<https://aibyhand.substack.com/p/11-can-you-calculate-self-attention>

### c) Multi-Head Attention

Instead of using just one attention function, the Transformer uses **multi-head attention**. We repeat the attention process  $h$  times (e.g.,  $h = 8$ ) with different learned projection matrices for  $W_Q, W_K, W_V$ .

The outputs from each head are concatenated:

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W_O$$

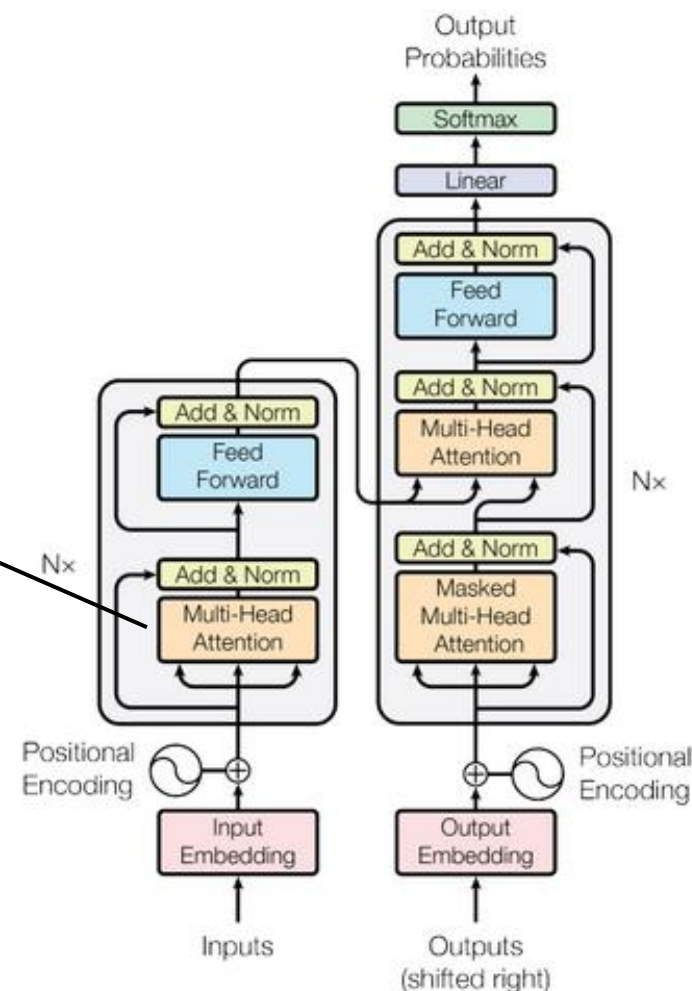
Where  $W_O \in \mathbb{R}^{h \cdot d_k \times d_{\text{model}}}$  is a projection matrix that reduces the concatenated output back to the original model dimension  $d_{\text{model}}$ .

### d) Add & Norm

The result of the multi-head attention is added to the input embeddings via a **residual connection**:

$$Z = \text{LayerNorm}(X' + \text{MultiHead}(Q, K, V))$$

This normalization step ensures stability during training by normalizing across the batch.





# 3. FFN (Encoder)

## Feed-Forward Network (FFN)

The FFN is a simple two-layer neural network applied independently to each position (token) in the sequence. The output of the self-attention mechanism (for each token) is passed through this network.

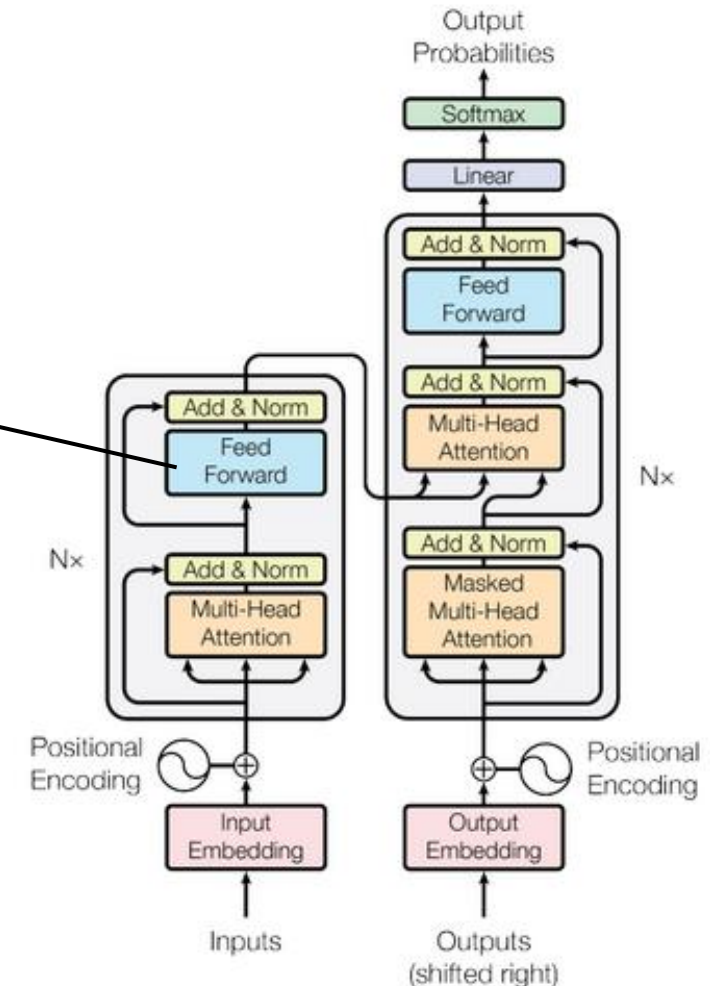
The FFN applies two linear transformations with a ReLU activation function in between.

The FFN Formula:

$$FFN(x) = \max(0, xW_1 + b_1)W_2 + b_2$$

Where:

- $x$ : The input, which is the output of the self-attention mechanism for a particular token (dimension  $d_{\text{model}}$ ).
- $W_1$ : A weight matrix of size  $d_{\text{model}} \times d_{\text{ff}}$ .
- $W_2$ : A weight matrix of size  $d_{\text{ff}} \times d_{\text{model}}$ .
- $b_1$ : A bias vector added after the first linear transformation (size  $d_{\text{ff}}$ ).
- $b_2$ : A bias vector added after the second linear transformation (size  $d_{\text{model}}$ ).
- $\text{ReLU}(\cdot)$ : The ReLU activation function applies the function  $\max(0, x)$  to introduce non-linearity.



# 4. Masked Multi-Head Attention (Decoder)

In the decoder, we generate the translation token by token. The decoder attends to the previously generated tokens, but it cannot attend to future tokens. This is enforced by a **mask**.

The process of **masked self-attention** is the same as regular self-attention, except that we apply a mask to prevent attending to future positions.

## a) Masking Mechanism

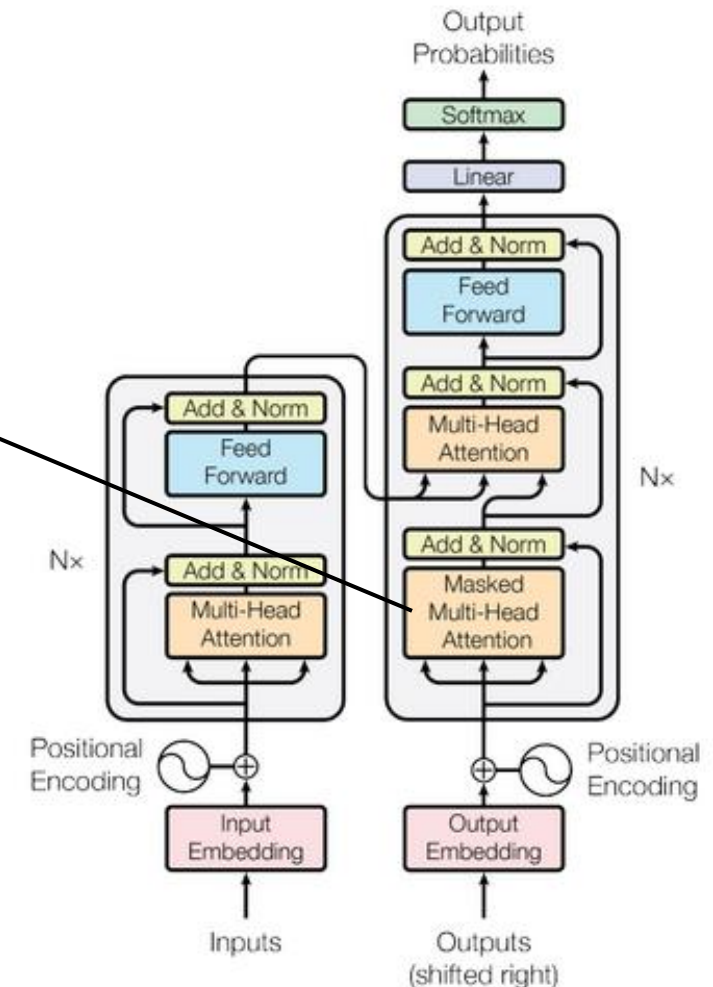
For each position  $t$  in the output, we prevent it from attending to future positions by setting the attention weights for those future positions to zero.

The masked attention score is calculated as:

$$\text{Masked Score} = \frac{QK^T}{\sqrt{d_k}} + \text{mask}$$

Where the mask has  $-\infty$  values for future positions and 0 for allowed positions.

The rest of the process (softmax, multiplication with values, and multi-head attention) follows the same procedure as in the encoder.



# 5. Encoder-Decoder Attention

In addition to masked self-attention, the decoder also attends to the encoder's output. This is done using the same multi-head attention mechanism, but this time, the **keys** and **values** come from the encoder, while the **queries** come from the decoder:

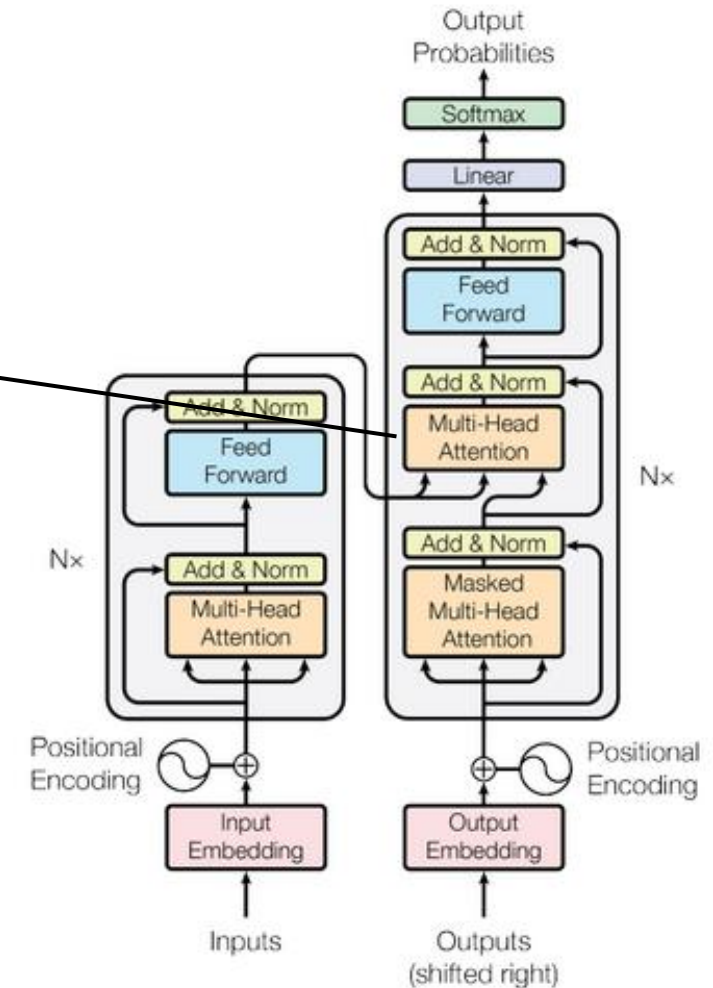
$$Q_{\text{decoder}} = Y'W_{Q_{\text{decoder}}}, \quad K_{\text{encoder}} = Z'W_{K_{\text{encoder}}}, \quad V_{\text{encoder}} = Z'W_{V_{\text{encoder}}}$$

The attention is computed in the same way as self-attention:

$$\text{Attention Weights}_{\text{enc-dec}} = \text{softmax} \left( \frac{Q_{\text{decoder}} K_{\text{encoder}}^T}{\sqrt{d_k}} \right)$$

$$\text{Attention Output}_{\text{enc-dec}} = \text{Attention Weights}_{\text{enc-dec}} \times V_{\text{encoder}}$$

This allows the decoder to attend to the context created by the encoder's representation of the input sentence.



# 6. Final Linear Layer and Softmax

The final output from the decoder is passed through a linear layer followed by a softmax to predict the next token.

## a) Linear Transformation

The output from the last decoder layer  $z_t$  is projected onto the vocabulary size  $V$  using a linear transformation:

$$y_t = W_z z_t + b_z$$

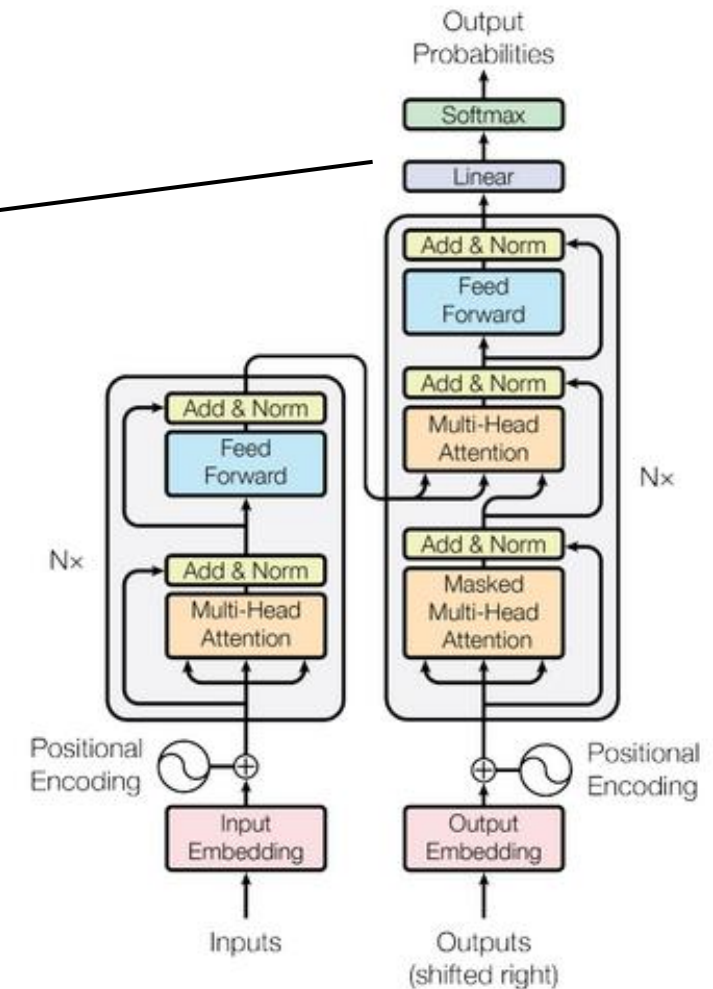
Where  $W_z \in \mathbb{R}^{d_{\text{model}} \times |V|}$  projects the decoder output to the size of the vocabulary.

## b) Softmax

Finally, a softmax is applied to obtain a probability distribution over all possible tokens:

$$\hat{y}_t = \text{softmax}(y_t)$$

The word with the highest probability is selected as the predicted token for the current position.



# Useful Links

- Vaswani, Ashish, et al. "Attention is all you need." Advances in neural information processing systems. 2017.
- <https://ai.googleblog.com/2017/08/transformer-novel-neural-network.html>
- <https://www.analyticsvidhya.com/blog/2019/06/understanding-transformers-nlp-state-of-the-art-models/>
- <https://jalammar.github.io/illustrated-transformer/>
- Sample Code Links:
  - <https://github.com/tensorflow/tensor2tensor/blob/master/README.md#walkthrough>
  - [https://colab.research.google.com/github/tensorflow/tensor2tensor/blob/master/tensor2tensor/notebooks/hello\\_t2t.ipynb#scrollTo=s19ucTii\\_wYb](https://colab.research.google.com/github/tensorflow/tensor2tensor/blob/master/tensor2tensor/notebooks/hello_t2t.ipynb#scrollTo=s19ucTii_wYb)